

Sharpness-Aware Minimization (SAM)


2025.02.13

KangJun Lee

Sharpness-Aware Minimization

Modern Neural Network 학습 과정에서…


Training Data와 Real Data를 같게 볼 수 있다고(i.i.d.) 가정하며 Training data의 Training Loss만 최소화시킴.



그러나, Training loss와 Real loss의 landscape가 다르며,
이로 인해 Training loss가 좋은 성능을 보장하지 못할 수 있다.

Sharpness-Aware Minimization

만약 Training distribution이 Test distribution에서 shift 된 형태라면,
Training Loss에서 Flat한 Minima이 Sharp한 Minima보다 더 좋은 Test Loss를 보일 것이다.



Idea) 그러면 Loss를 최소화할 뿐만 아니라 Flat한 minima를 찾으면 높은 Generalization ability를 보이지 않을까?

Sharpness-Aware Minimization

(Flat한지 인지하면서)

(Loss를 Minimizing하는 Optimizer)

SHARPNESS-AWARE MINIMIZATION FOR EFFICIENTLY IMPROVING GENERALIZATION

Pierre Foret *

Google Research

pierre.pforet@gmail.com

Ariel Kleiner

Google Research

akleiner@gmail.com

Hossein Mobahi

Google Research

hmobahi@google.com

Idea) Flat한 minima를 찾는 Optimizer는 Generalization ability를 향상시켜줄 것이다.

Sharpness-Aware Minimization

Theorem (stated informally) 1. *For any $\rho > 0$, with high probability over training set \mathcal{S} generated from distribution \mathcal{D} ,*

$$L_{\mathcal{D}}(\mathbf{w}) \leq \max_{\|\boldsymbol{\epsilon}\|_2 \leq \rho} L_{\mathcal{S}}(\mathbf{w} + \boldsymbol{\epsilon}) + h(\|\mathbf{w}\|_2^2 / \rho^2),$$

where $h : \mathbb{R}_+ \rightarrow \mathbb{R}_+$ is a strictly increasing function (under some technical conditions on $L_{\mathcal{D}}(\mathbf{w})$).

To make explicit our sharpness term, we can rewrite the right hand side of the inequality above as

$$[\max_{\|\boldsymbol{\epsilon}\|_2 \leq \rho} L_{\mathcal{S}}(\mathbf{w} + \boldsymbol{\epsilon}) - L_{\mathcal{S}}(\mathbf{w})] + L_{\mathcal{S}}(\mathbf{w}) + h(\|\mathbf{w}\|_2^2 / \rho^2).$$

-> PAC-Bayesian Bound를 활용해서 증명. (Skip the Proof)

Q. PAC-Bayesian Bound?

PAC(Probabilistic Approximately Correct) Bound

Meaning : 높은 확률로 (high Probability), true and empirical risk의 차이를 근사한다. (Approximately).
(Real and Sample data loss)

Example) Hoeffding Inequality

Let Z_1, \dots, Z_N be i.i.d. random variables bounded in $[0, 1]$. Then for all $\epsilon > 0$,

$$\mathbb{P} \left[\left| \frac{1}{N} \sum_{n=1}^N Z_n - \mathbb{E}[Z_n] \right| > \epsilon \right] \leq 2 \exp(-2N\epsilon^2).$$

By writing $\delta = 2\exp(-2N\epsilon^2)$, we get, with probability $1 - \delta$, High Probability

$$\left| \frac{1}{N} \sum_{n=1}^N Z_n - \mathbb{E}[Z_n] \right| \leq \sqrt{\frac{1}{2N} \log \frac{2}{\delta}}$$

Approximately

PAC(Probabilistic Approximately Correct) Bayesian-Bound

PAC 개념으로부터, Bayesian 개념을 차용하여 PAC Bayesian-Bound을 유도할 수 있다.

Recall) Bayesian Theory

$$P(A|B) = \frac{P(B|A)P(A)}{P(B)}$$

↑ ↓
Posterior Probability Predictor Prior Probability
Likelihood Class Prior Probability

-> 기존에 주어진 확률(Prior Probability, Likelihood)를 통해
사후 확률(Posterior Probability)을 구할 수 있다.

PAC(Probabilistic Approximately Correct) Bayesian-Bound

여기에서 Prior, Posterior의 개념을 차용해서,

prior P 를 Sample data와 관계없는 hypothesis(predictor)로 정의할 수 있다.

posterior $Q(S) = Q$ 를 Sample data에 종속된 hypothesis(predictor)로 정의할 수 있다.

여기에서, Prior와 Posterior는 (Bayesian과 다르게) 사건이 아닌 Probability Distribution이다.

그러면 모든 hypothesis(predictor), h 에 대해 real, empirical risk를 정의할 수 있다.

$$R_Q = \mathbb{E}_{h \sim Q}[R(h)] = \mathbb{E}_{h \sim Q} \left[\mathbb{E}_{(x,y) \sim D}[l((x,y), h)] \right]$$

$$r_{S,Q} = \mathbb{E}_{h \sim Q}[r_s(h)] = \mathbb{E}_{h \sim Q} \left[\frac{1}{N} \sum_{(x,y) \in S} l((x,y), h) \right]$$

PAC(Probabilistic Approximately Correct) Bayesian-Bound

이를 이용하여, PAC-Bayesian을 정의할 수 있다.

$$R_Q \leq^{\text{Probabilistic}} r_{S,Q} + \underline{f(Q, P, N, \delta)}^{\text{Approximately}}$$

Ex) Mcallester's Thm (1999)

For any $\ell \in \{0, 1\}$, D, \mathcal{H} and P a probability measure supported on \mathcal{H} ,
for $N \geq 8$,

$$R_Q \stackrel{\text{Probabilistic}}{\leq} r_{S,Q} + \sqrt{\frac{\mathcal{D}_{\text{KL}}[Q||P] + \log \sqrt{N} + \log \frac{2}{\delta}}{2N}} \quad (1)$$

for all Q probability measures supported on \mathcal{H} . Approximately

(Mcallester's Thm이 실제 증명에 사용됨.)

Sharpness-Aware Minimization

Theorem (stated informally) 1. *For any $\rho > 0$, with high probability over training set \mathcal{S} generated from distribution \mathcal{D} ,*

$$L_{\mathcal{D}}(\mathbf{w}) \leq \max_{\|\epsilon\|_2 \leq \rho} L_{\mathcal{S}}(\mathbf{w} + \epsilon) + h(\|\mathbf{w}\|_2^2 / \rho^2),$$

where $h : \mathbb{R}_+ \rightarrow \mathbb{R}_+$ is a strictly increasing function (under some technical conditions on $L_{\mathcal{D}}(\mathbf{w})$).

To make explicit our sharpness term, we can rewrite the right hand side of the inequality above as

$$\underbrace{[\max_{\|\epsilon\|_2 \leq \rho} L_{\mathcal{S}}(\mathbf{w} + \epsilon) - L_{\mathcal{S}}(\mathbf{w})]}_{\text{Sharpness Term}} + \underbrace{L_{\mathcal{S}}(\mathbf{w}) + h(\|\mathbf{w}\|_2^2 / \rho^2)}_{\text{Loss & Regularization}}.$$

Sharpness Term



Loss & Regularization

Motivation : loss와 sharpness를 동시에 최소화 하는 것이 실제 데이터에서 좋은 성능을 갖게 한다.

-> Good Generalization ability

Sharpness-Aware Minimization

$$\frac{[\max_{\|\epsilon\|_2 \leq \rho} L_S(\mathbf{w} + \epsilon) - L_S(\mathbf{w})] + L_S(\mathbf{w}) + h(\|\mathbf{w}\|_2^2 / \rho^2)}{\text{Sharpness Term} \quad \text{Loss & Regularization}}$$



Motivation : loss와 sharpness를 동시에 최소화하자!

Sharpness-Aware Minimization

결론적으로 둘 다 줄이기 위해서는, 아래 식을 최소화 해야한다.

$$\min_w L_s^{SAM}(w) + \lambda \|w\|_2^2$$

where

$$L_s^{SAM}(w) \triangleq \max_{\|\epsilon\|_p \leq \rho} L_s(w + \epsilon)$$

$\max_{\|\epsilon\|_p \leq \rho} L_s(w + \epsilon)$ 를 구하기 위해서, $\max_{\|\epsilon\|_p \leq \rho} L_s(w + \epsilon) = L_s(w + \hat{\epsilon})$ 만족하는 $\hat{\epsilon}(w)$ 를 구해야한다.

$$\hat{\epsilon}(w) = \operatorname{argmax}_{\|\epsilon\|_p \leq \rho} L_s(w + \epsilon) \approx \operatorname{argmax}_{\|\epsilon\|_p \leq \rho} L_s(w) + \epsilon^T \nabla_w L_s(w) = \operatorname{argmax}_{\|\epsilon\|_p \leq \rho} \epsilon^T \nabla_w L_s(w)$$

$$\rightarrow \hat{\epsilon}(w) = \rho \frac{\nabla_w L_s(w)}{\left\| \nabla_w L_s(w) \right\|_2^2}$$

Gradient Calculation

Sharpness-Aware Minimization

그러면,

$$\nabla_w L_s^{SAM}(w) \approx \nabla_w L_s(w + \hat{\epsilon}) = \frac{d(w + \hat{\epsilon})}{dw} \nabla_w L_s(w) \Big|_{w+\hat{\epsilon}}$$

$$\frac{d(w + \hat{\epsilon})}{dw} \nabla_w L_s(w) \Big|_{w+\hat{\epsilon}} = \nabla_w L_s(w) \Big|_{w+\hat{\epsilon}} + \frac{d\hat{\epsilon}}{dw} \nabla_w L_s(w) \Big|_{w+\hat{\epsilon}}$$

To accelerate Computation

$$\therefore \nabla_w L_s^{SAM}(w) \approx \nabla_w L_s(w) \Big|_{w+\hat{\epsilon}}$$

Gradient Calculation

전체 과정에서 Gradient가 2번 계산된다. \rightarrow More Computation Time

Sharpness-Aware Minimization

Algorithm)

Input: Training set $\mathcal{S} \triangleq \cup_{i=1}^n \{(\mathbf{x}_i, \mathbf{y}_i)\}$, Loss function $l : \mathcal{W} \times \mathcal{X} \times \mathcal{Y} \rightarrow \mathbb{R}_+$, Batch size b , Step size $\eta > 0$, Neighborhood size $\rho > 0$.

Output: Model trained with SAM

Initialize weights \mathbf{w}_0 , $t = 0$;

while *not converged* **do**

- Sample batch $\mathcal{B} = \{(\mathbf{x}_1, \mathbf{y}_1), \dots, (\mathbf{x}_b, \mathbf{y}_b)\}$;
- Compute gradient $\nabla_{\mathbf{w}} L_{\mathcal{B}}(\mathbf{w})$ of the batch's training loss;
- Compute $\hat{\epsilon}(\mathbf{w})$ per equation 2;
- Compute gradient approximation for the SAM objective (equation 3): $\mathbf{g} = \nabla_{\mathbf{w}} L_{\mathcal{B}}(\mathbf{w})|_{\mathbf{w} + \hat{\epsilon}(\mathbf{w})}$;
- Update weights: $\mathbf{w}_{t+1} = \mathbf{w}_t - \eta \mathbf{g}$;
- $t = t + 1$;

end

return \mathbf{w}_t

Algorithm 1: SAM algorithm




Figure 2: Schematic of the SAM parameter update.

Experimental Result (Image Classification)

Experimental-Setup)

| Model | Augmentation |
|-------------------------|--------------|
| WRN-28-10 (200 epochs) | Basic |
| WRN-28-10 (200 epochs) | Cutout |
| WRN-28-10 (200 epochs) | AA |
| WRN-28-10 (1800 epochs) | Basic |
| WRN-28-10 (1800 epochs) | Cutout |
| WRN-28-10 (1800 epochs) | AA |
| Shake-Shake (26 2x96d) | Basic |
| Shake-Shake (26 2x96d) | Cutout |
| Shake-Shake (26 2x96d) | AA |
| PyramidNet | Basic |
| PyramidNet | Cutout |
| PyramidNet | AA |
| PyramidNet+ShakeDrop | Basic |
| PyramidNet+ShakeDrop | Cutout |
| PyramidNet+ShakeDrop | AA |

다음의 NN들을 대상으로 SAM과 SGD의 정확도를 비교함.

SAM에서 ρ 값은 {0.01, 0.02, 0.05, 0.1, 0.2, 0.5}의 값 중 가장 성능이 좋은 값으로 설정함. (Grid Search)

또한, SAM 학습 과정에서 8대의 Nvidia V100 GPU를 통해 병렬로 계산하여 각각의 data에서 계산한 ϵ 값의 평균을 구하여 update함. (뒤에 나올 m-Sharpness 방식과 유사함.)

모든 데이터에 대해서 basic data augmentations (horizontal flip, padding by four pixels, and random crop) 이 적용됨.

Experimental Result (Image Classification)

| Model | Augmentation | CIFAR-10 | | CIFAR-100 | |
|-------------------------|--------------|--------------------------------|-------------------------|-----------------------------|----------------------|
| | | SAM | SGD | SAM | SGD |
| WRN-28-10 (200 epochs) | Basic | 2.7 _{±0.1} | 3.5 _{±0.1} | 16.5 _{±0.2} | 18.8 _{±0.2} |
| WRN-28-10 (200 epochs) | Cutout | 2.3 _{±0.1} | 2.6 _{±0.1} | 14.9 _{±0.2} | 16.9 _{±0.1} |
| WRN-28-10 (200 epochs) | AA | 2.1 _{±<0.1} | 2.3 _{±0.1} | 13.6 _{±0.2} | 15.8 _{±0.2} |
| WRN-28-10 (1800 epochs) | Basic | 2.4 _{±0.1} | 3.5 _{±0.1} | 16.3 _{±0.2} | 19.1 _{±0.1} |
| WRN-28-10 (1800 epochs) | Cutout | 2.1 _{±0.1} | 2.7 _{±0.1} | 14.0 _{±0.1} | 17.4 _{±0.1} |
| WRN-28-10 (1800 epochs) | AA | 1.6 _{±0.1} | 2.2 _{±<0.1} | 12.8 _{±0.2} | 16.1 _{±0.2} |
| Shake-Shake (26 2x96d) | Basic | 2.3 _{±<0.1} | 2.7 _{±0.1} | 15.1 _{±0.1} | 17.0 _{±0.1} |
| Shake-Shake (26 2x96d) | Cutout | 2.0 _{±<0.1} | 2.3 _{±0.1} | 14.2 _{±0.2} | 15.7 _{±0.2} |
| Shake-Shake (26 2x96d) | AA | 1.6 _{±<0.1} | 1.9 _{±0.1} | 12.8 _{±0.1} | 14.1 _{±0.2} |
| PyramidNet | Basic | 2.7 _{±0.1} | 4.0 _{±0.1} | 14.6 _{±0.4} | 19.7 _{±0.3} |
| PyramidNet | Cutout | 1.9 _{±0.1} | 2.5 _{±0.1} | 12.6 _{±0.2} | 16.4 _{±0.1} |
| PyramidNet | AA | 1.6 _{±0.1} | 1.9 _{±0.1} | 11.6 _{±0.1} | 14.6 _{±0.1} |
| PyramidNet+ShakeDrop | Basic | 2.1 _{±0.1} | 2.5 _{±0.1} | 13.3 _{±0.2} | 14.5 _{±0.1} |
| PyramidNet+ShakeDrop | Cutout | 1.6 _{±<0.1} | 1.9 _{±0.1} | 11.3 _{±0.1} | 11.8 _{±0.2} |
| PyramidNet+ShakeDrop | AA | 1.4 _{±<0.1} | 1.6 _{±<0.1} | 10.3 _{±0.1} | 10.6 _{±0.1} |

(Error-rate)

SGD와 비교했을 때, 모든 Data Augmentation 기법과, 다른 Network에 대해서도 유의미하게 성능 향상을 나타냄.

(여기에서, SAM의 1 epoch당 시간이 2배나 걸리기 때문에 SAM에서는 200, SGD에서는 epoch=400으로 학습함.)

Experimental Result (Image Classification)

| Model | Epoch | SAM | | Standard Training (No SAM) | |
|------------|-------|-------------------------|-----------------|----------------------------|-----------------|
| | | Top-1 | Top-5 | Top-1 | Top-5 |
| ResNet-50 | 100 | 22.5 \pm 0.1 | 6.28 \pm 0.08 | 22.9 \pm 0.1 | 6.62 \pm 0.11 |
| | 200 | 21.4 \pm 0.1 | 5.82 \pm 0.03 | 22.3 \pm 0.1 | 6.37 \pm 0.04 |
| | 400 | 20.9 \pm 0.1 | 5.51 \pm 0.03 | 22.3 \pm 0.1 | 6.40 \pm 0.06 |
| ResNet-101 | 100 | 20.2 \pm 0.1 | 5.12 \pm 0.03 | 21.2 \pm 0.1 | 5.66 \pm 0.05 |
| | 200 | 19.4 \pm 0.1 | 4.76 \pm 0.03 | 20.9 \pm 0.1 | 5.66 \pm 0.04 |
| | 400 | 19.0 \pm <0.01 | 4.65 \pm 0.05 | 22.3 \pm 0.1 | 6.41 \pm 0.06 |
| ResNet-152 | 100 | 19.2 \pm <0.01 | 4.69 \pm 0.04 | 20.4 \pm <0.0 | 5.39 \pm 0.06 |
| | 200 | 18.5 \pm 0.1 | 4.37 \pm 0.03 | 20.3 \pm 0.2 | 5.39 \pm 0.07 |
| | 400 | 18.4 \pm <0.01 | 4.35 \pm 0.04 | 20.9 \pm <0.0 | 5.84 \pm 0.07 |

Table 2: Test error rates for ResNets trained on ImageNet, with and without SAM.

(Without SAM) / SGD with cosine learning-schedule, momentum with 0.9, weight-decay = 0.005.

앞의 예시보다 비교적 무거운 모델인 ResNet에 대해서도 SAM이 더 좋은 성능을 보임.

심지어, SAM은 epoch가 많아지더라도 overfitting 현상이 일어나지 않은 것을 볼 수 있음. (Regularization Term)

Experimental Result (Fine-Tuning)

| Dataset | EffNet-b7 + SAM | EffNet-b7 | Prev. SOTA (ImageNet only) | EffNet-L2 + SAM | EffNet-L2 | Prev. SOTA |
|------------------|-------------------------|------------------|-------------------------------|-------------------------|------------------|------------------|
| FGVC_Aircraft | 6.80 ± 0.06 | 8.15 ± 0.08 | 5.3 (TBMSL-Net) | 4.82 ± 0.08 | 5.80 ± 0.1 | 5.3 (TBMSL-Net) |
| Flowers | 0.63 ± 0.02 | 1.16 ± 0.05 | 0.7 (BiT-M) | 0.35 ± 0.01 | 0.40 ± 0.02 | 0.37 (EffNet) |
| Oxford_IIIT_Pets | 3.97 ± 0.04 | 4.24 ± 0.09 | 4.1 (Gpipe) | 2.90 ± 0.04 | 3.08 ± 0.04 | 4.1 (Gpipe) |
| Stanford_Cars | 5.18 ± 0.02 | 5.94 ± 0.06 | 5.0 (TBMSL-Net) | 4.04 ± 0.03 | 4.93 ± 0.04 | 3.8 (DAT) |
| CIFAR-10 | 0.88 ± 0.02 | 0.95 ± 0.03 | 1 (Gpipe) | 0.30 ± 0.01 | 0.34 ± 0.02 | 0.63 (BiT-L) |
| CIFAR-100 | 7.44 ± 0.06 | 7.68 ± 0.06 | 7.83 (BiT-M) | 3.92 ± 0.06 | 4.07 ± 0.08 | 6.49 (BiT-L) |
| Birdsnap | 13.64 ± 0.15 | 14.30 ± 0.18 | 15.7 (EffNet) | 9.93 ± 0.15 | 10.31 ± 0.15 | 14.5 (DAT) |
| Food101 | 7.02 ± 0.02 | 7.17 ± 0.03 | 7.0 (Gpipe) | 3.82 ± 0.01 | 3.97 ± 0.03 | 4.7 (DAT) |
| ImageNet | 15.14 ± 0.03 | 15.3 | 14.2 (KDforAA) | 11.39 ± 0.02 | 11.8 | 11.45 (ViT) |


(Error-rate)

Fine-Tuning의 경우에서도 SAM이 기존에 쓰였던 방법에 비해 이점을 보임.

(여기에서, EffNet-b7과 EffNet-L2는 각각 RandAugment와 NoisyStudent 기법으로 학습된 모델의 checkpoint를 가져와서 Fine-Tuning 했다고 함.)

Experimental Result (Self-Experiment)

Experimental-Setup)



학습은 개인 노트북으로 진행.

(RTX 1650Ti)

Task : CIFAR10 Classification

Github에 있는 코드를 이용하여 학습.


Data augmentation 적용, momentum = 0.9,
Batch_size = 64

Weight_decay = 0.0005

Step_LR 적용

$\rho = 0.5$

Experimental Result (Self-Experiment)



| Optimizer | Loss | Accuracy(%) |
|------------|---------|-------------|
| SGD | 0.60742 | 74.42 |
| SAM_single | 0.59017 | 75.74 |
| SAM_multi | 0.58964 | 76.27 |

앞의 예시와 같이 가벼운 모델에서도 SAM이 유의미한 성능 향상을 보임.

Multi-Step(double) SAM 같은 경우 SAM과 비교하여 약간의 성능 향상을 보임.

특히, 학습 후반부에 Multi-Step SAM이 기존 SAM에 비해 더 좋은 성능을 보임.

Experimental Result (Self-Experiment)

Q. Multi-Step SAM도 Second-order term을 무시할 수 있을까?

$$\nabla L_s^{SAM}(w) \approx \nabla_w L_s(w) \Big|_{w+\sum_{n=1}^m \hat{\epsilon}_n} + \frac{d(\sum_{n=1}^m \hat{\epsilon}_n)}{dw} \nabla_w L_s(w) \Big|_{w+\sum_{n=1}^m \hat{\epsilon}_n}$$

$\frac{d(\sum_{n=1}^m \hat{\epsilon}_n)}{dw} \nabla_w L_s(w) \Big|_{w+\sum_{n=1}^m \hat{\epsilon}_n}$ 부분도 똑같이 무시될 수 있는데, 이 부분이

Single-Step SAM과 다르게 나타날 수 있지 않을까?


그러나, $\left\| \sum_{n=1}^m \hat{\epsilon}_n \right\| \leq \rho$, 즉 permutation은 항상 $B_\rho(w)$ 안에 존재한다.

Single-Step에서 $\frac{d\hat{\epsilon}}{dw} \nabla_w L_s(w) \Big|_{w+\hat{\epsilon}}$ 가 무시될 수 있고, $\|\hat{\epsilon}\| \leq \rho$ 임을 고려하면,

같은 이유로 $\frac{d(\sum_{n=1}^m \hat{\epsilon}_n)}{dw} \nabla_w L_s(w) \Big|_{w+\sum_{n=1}^m \hat{\epsilon}_n}$ 도 무시될 수 있다.

m-step SAM

$$\begin{aligned} \nabla \hat{\Lambda}^{SAM}(w) &= \nabla \hat{\Lambda}(w + \sum_{n=1}^m \hat{\epsilon}_n) \approx \frac{d(\sum_{n=1}^m \hat{\epsilon}_n)}{dw} \nabla \hat{\Lambda}(w) \Big|_{w + \sum_{n=1}^m \hat{\epsilon}_n} \\ \Rightarrow \text{ consider } &\frac{d(\sum_{n=1}^m \hat{\epsilon}_n)}{dw} \nabla \hat{\Lambda}(w) \Big|_{w + \sum_{n=1}^m \hat{\epsilon}_n}. \quad (\text{here, } \|\hat{\epsilon}_n\| = \frac{\rho}{m}) \end{aligned}$$



but $\left\| \sum_{n=1}^m \hat{\epsilon}_n \right\| \leq \rho \Rightarrow$ Considering we can ignore $\|\hat{\epsilon}\| \leq \rho$,

$\frac{d\hat{\epsilon}}{dw} \nabla \hat{\Lambda}(w) \Big|_{w+\hat{\epsilon}}$, $\frac{d(\sum_{n=1}^m \hat{\epsilon}_n)}{dw} \nabla \hat{\Lambda}(w) \Big|_{w + \sum_{n=1}^m \hat{\epsilon}_n}$. can be ignored.

Experimental Result (Self-Experiment)

Q. Second-Order Term이 무시될 수 있는 이유?




Figure 4: Training and test error for the first and second order version of the algorithm.




Figure 5: Cosine similarity between the first and second order updates.

논문에서는 실험적으로 보여줌.

학습 중반까지, First order update SAM의 학습 방향이
Second order update SAM과 방향이 거의 같음을 볼 수 있다.

중반 이후에는 방향이 달라지지만,
second-order을 생략한 SAM 모델이 더 낮은 train error를 보임.

Experimental Result (m-Sharpness)

Input: Training set $\mathcal{S} \triangleq \cup_{i=1}^n \{(\mathbf{x}_i, \mathbf{y}_i)\}$, Loss function $l : \mathcal{W} \times \mathcal{X} \times \mathcal{Y} \rightarrow \mathbb{R}_+$, Batch size b , Step size $\eta > 0$, Neighborhood size $\rho > 0$.

Output: Model trained with SAM

Initialize weights \mathbf{w}_0 , $t = 0$;

while *not converged* **do**

- Sample batch $\mathcal{B} = \{(\mathbf{x}_1, \mathbf{y}_1), \dots, (\mathbf{x}_b, \mathbf{y}_b)\}$;
- Compute gradient $\nabla_{\mathbf{w}} L_{\mathcal{B}}(\mathbf{w})$ of the batch's training loss;
- Compute $\hat{\epsilon}(\mathbf{w})$ per equation 2;
- Compute gradient approximation for the SAM objective (equation 3): $\mathbf{g} = \nabla_{\mathbf{w}} L_{\mathcal{B}}(\mathbf{w})|_{\mathbf{w} + \hat{\epsilon}(\mathbf{w})}$;
- Update weights: $\mathbf{w}_{t+1} = \mathbf{w}_t - \eta \mathbf{g}$;
- $t = t + 1$;

end

return \mathbf{w}_t

Algorithm 1: SAM algorithm


기존 알고리즘)

Batch를 나누어서 각 Batch마다 ϵ 을 구한 뒤 weight를 update.

m-Sharpness)


각 배치마다 독립적으로 계산한 뒤 각 배치마다의 ϵ 의 평균 perturbation을 구하여 그 값을 바탕으로 weight를 update.

Experimental Result (m-Sharpness)



(left) Test error as a function of ρ for different values of m .

(right) Predictive power of m-sharpness for the generalization gap, for different values of m (higher means the sharpness measure is more correlated with actual generalization gap).




배치 크기, m 의 값이 작아질수록 더 좋은 generalization ability를 보여줌.

이 사실은 대규모 학습에서 필요한 병렬학습에도 적합하다는 사실을 보여줌.

또한, m 의 값이 작을수록 sharpness-term이 generalization에 영향을 더 많이 끼친다는 것을 보여줌.

Experimental Result (Hessian-Spectra)



Evolution of the spectrum of the Hessian during training of a model with standard SGD (lefthand column) or SAM (righthand column).

마지막으로, Epoch에 따라 Hessian Matrix의 eigenvalue를 비교함.

Hessian의 Eigenvalue가 작을수록 평탄하다고 볼 수 있는데, SGD와 비교했을 때 SAM의 eigenvalue가 월등히 작음을 알 수 있음.

그렇기에, loss 뿐만 아니라 Flatness까지 잘 학습되고 있다는 것을 알 수 있음.

Thank you!